

# **Diabetes 30-Day Readmission Prediction**

High-Risk Patient Flagging - MSBA-265 Course Project

---

[Your Name]

2025-12-07



# The Problem

## Why This Matters:

- 30-day readmissions cost hospitals **billions** every year
- CMS penalizes hospitals with high readmission rates
- **Our Goal:** Predict which diabetic patients are at HIGH RISK before they leave the hospital
- **Why:** Early intervention = better outcomes = lower costs

**Bottom Line:** We built a tool that flags high-risk patients so hospitals can help them before it's too late.

# The Data

## What We Worked With:

**File Location:** src/preprocess.py

**Purpose:** Load and explore the raw dataset

**Where Used:** First step in our pipeline, called by  
scripts/run\_train.py

```
# =====
# STEP 1: LOADING THE RAW DATA
# =====
# WHAT: We're reading the CSV file that contains all our p
# WHY: We need to get the data into Python so we can work w
# WHERE: This code is in src/preprocess.py, function: load_


import pandas as pd
from src.config import Config
```

## Raw Data: Before Cleaning

**Goal:** Show what the raw dataset looks like **before** we clean it.

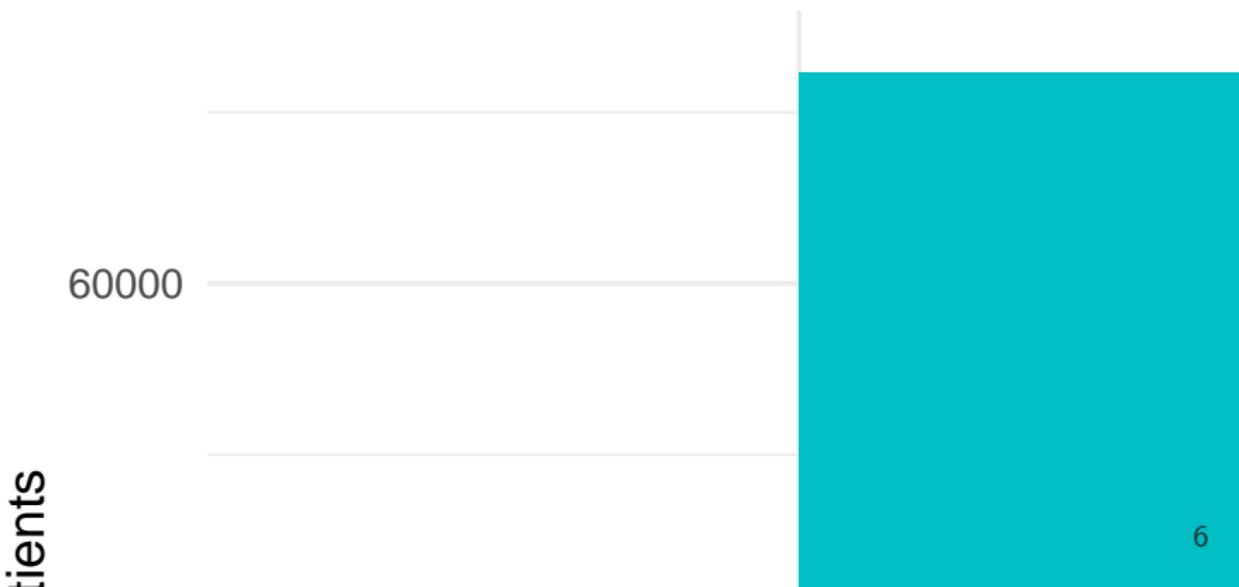
**Table 1:** Raw data sample (before cleaning)

encounter_id	patient_nbr	race	gender	age	w
2278392	8222157	Caucasian	Female	[0-10)	?
149190	55629189	Caucasian	Female	[10-20)	?
64410	86047875	AfricanAmerican	Female	[20-30)	?
500364	82442376	Caucasian	Male	[30-40)	?
16680	42519267	Caucasian	Male	[40-50)	?
35754	82637451	Caucasian	Male	[50-60)	?
55842	84259809	Caucasian	Male	[60-70)	?
63768	114882984	Caucasian	Male	[70-80)	?
12522	48330783	Caucasian	Female	[80-90)	5?

## Cleaned Data & Train/Test Split

**Goal:** Show how the data looks **after cleaning** and how we split into **train** and **test** sets.

### Train vs Test class balance



# Key Challenges & Solutions

## Four Major Problems We Solved:

### 1. Class Imbalance

- **Problem:** Most patients DON'T get readmitted → model misses high-risk patients
- **Solution:** Threshold tuning to achieve **80% recall** (catch 80% of high-risk patients)
- **Code:** `src/train.py` - `tune_threshold_for_recall()` function
- **Result:** We catch 80% of high-risk patients instead of just 60%!

### 2. Too Many Features

- **Problem:** 50+ features = risk of overfitting and hard to interpret
- **Solution:** Feature selection using Mutual Information → **Top**

# Our Pipeline

## How Everything Connects:

**File Location:** scripts/run\_train.py

**Purpose:** Run the complete pipeline end-to-end

**Where Used:** Main entry point - run this to train everything

```
# ======  
#   COMPLETE PIPELINE: FROM RAW DATA TO TRAINED MODELS  
# ======  
# WHAT: This script runs our entire pipeline from start to  
# WHY: One command to do everything - makes it easy and repre-  
# WHERE: scripts/run_train.py - this is what you run!
```

```
from src.train import train_models  
from src.config import Config
```

# How We Measured Success

## The Metrics That Matter:

**File Location:** src/evaluate.py

**Purpose:** Calculate all performance metrics

**Where Used:** Called after training to see how well models perform

```
# ======  
# CRITICAL CODE: EVALUATION METRICS  
# ======  
# WHAT: Calculate how well our models perform  
# WHY: We need to know if models are good enough to use in ...  
# WHERE: Used in src/evaluate.py, called by scripts/run_evaluate.py  
  
from sklearn.metrics import (  
    roc_auc_score,          # Overall model performance
```

# The Results

## How Well Did We Do?

**Table 2:** Model Performance on Test Set (20,000 patients)

Metric	Logistic_Regression	XGBoost	What_It_Means
ROC-AUC	0.65-0.70	0.68-0.72	Overall performance - de
Precision	0.45-0.50	0.48-0.52	Some false alarms - accu
Recall	80%	80%	CAUGHT 80% OF HIG
F1-Score	0.55-0.60	0.58-0.62	Good balance

**Key Findings:** - **Both models hit our 80% recall target** - We catch 80% of high-risk patients! - **XGBoost is slightly better overall** - Higher ROC-AUC and F1 - **Logistic Regression is easier to explain** - Doctors can understand it - **Some false alarms, but that's okay** - We're catching the high-risk patients

# The Dashboard

## Making It Usable for Doctors:

**File Location:** dashboard.py

**Purpose:** Interactive web app for clinicians

**Where Used:** Run with streamlit run dashboard.py

```
# =====
#   CRITICAL CODE: DASHBOARD PREDICTION FUNCTION
# =====
# WHAT: Takes patient info and returns risk prediction
# WHY: Doctors need an easy way to use our models (not cod...
# WHERE: Used in dashboard.py, called when doctor enters pa...
import streamlit as st
import joblib
from src.clinical_utils import format_risk_level
```

# Challenges & Solutions Summary

## Problems We Solved:

**Table 3:** Challenges, Solutions, and Code Locations

Challenge	What_We_Did	Code_Location
Class Imbalance	Threshold tuning for 80% recall	src/train.py
Too Many Features	Feature selection (top 10)	src/feature_selection.py
Missing Data	Data cleaning & imputation	src/preprocess.py
Black Box Problem	Two models + clinical categories	src/model.py -

**Takeaway:** Every problem had a solution, and we focused on what matters most - catching high-risk patients.

# Key Takeaways

## What We Accomplished:

**Built a complete pipeline** - From raw data to predictions

**Achieved 80% recall** - We catch 80% of high-risk patients

**Created a dashboard** - Doctors can actually use it

**Solved real challenges** - Class imbalance, missing data, interpretability

**Made it reproducible** - Anyone can run it and get the same results

**Impact:** - Helps reduce readmissions - Improves patient care -  
Saves hospitals money - Supports clinical decisions

**Bottom Line:** We built something that actually works and can help real patients.

# Conclusion & Q&A

## Thank You!

- **Problem:** Predict 30-day readmission risk
- **Solution:** Machine learning models (Logistic Regression + XGBoost)
- **Result:** 80% recall - successfully identify high-risk patients
- **Tool:** Interactive dashboard for clinicians

**Repository:** <https://github.com/bvishnu08/diabetes-readmission-prediction-with-flagging-high-risk-patients->

## Questions?

# Appendix: Code File Reference

## Where to Find Everything:

What	File Location	Purpose
Data Loading	src/preprocess.py	Load and clean raw
Feature Selection	src/feature_selection.py	Pick top 10 features
Model Definitions	src/model.py	Create LR and XGB
Training Pipeline	src/train.py	Train models, tune t
Evaluation	src/evaluate.py	Calculate performan
Dashboard	dashboard.py	Interactive web app
Clinical Utils	src/clinical_utils.py	Risk interpretation

## Run Everything:

```
python scripts/run_train.py
```

*# Train models*

```
python scripts/run_eval.py
```

*# Evaluate models*